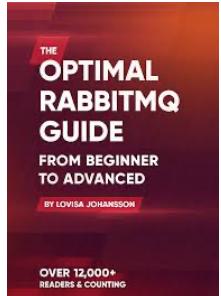
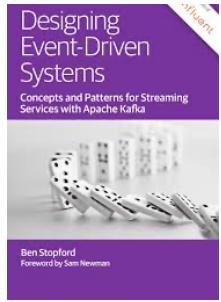


Event-driven systems

UA.DETI.IES - 2019/20

Resources & Credits



- ❖ Designing Event-Driven Systems
Ben Stopford, O'Reilly

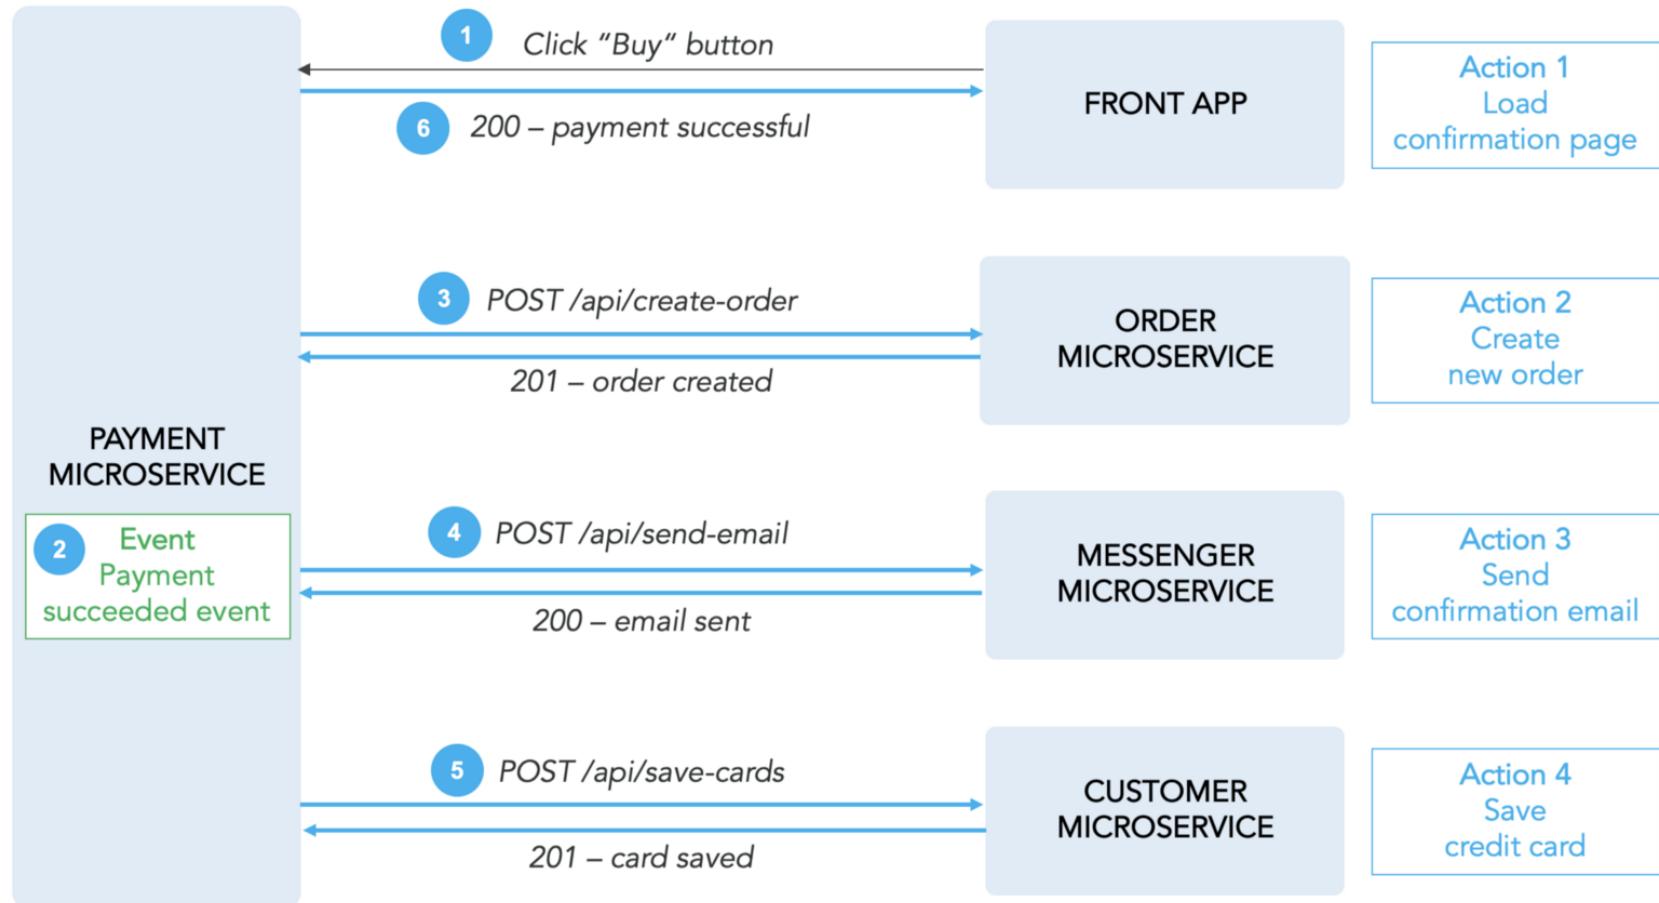
- ❖ The Optimal RabbitMQ Guide, Lovisa Johansson, CloudAMQP

- ❖ Spring Cloud Data Flow
 - <https://dataflow.spring.io>

Scaling with Microservices

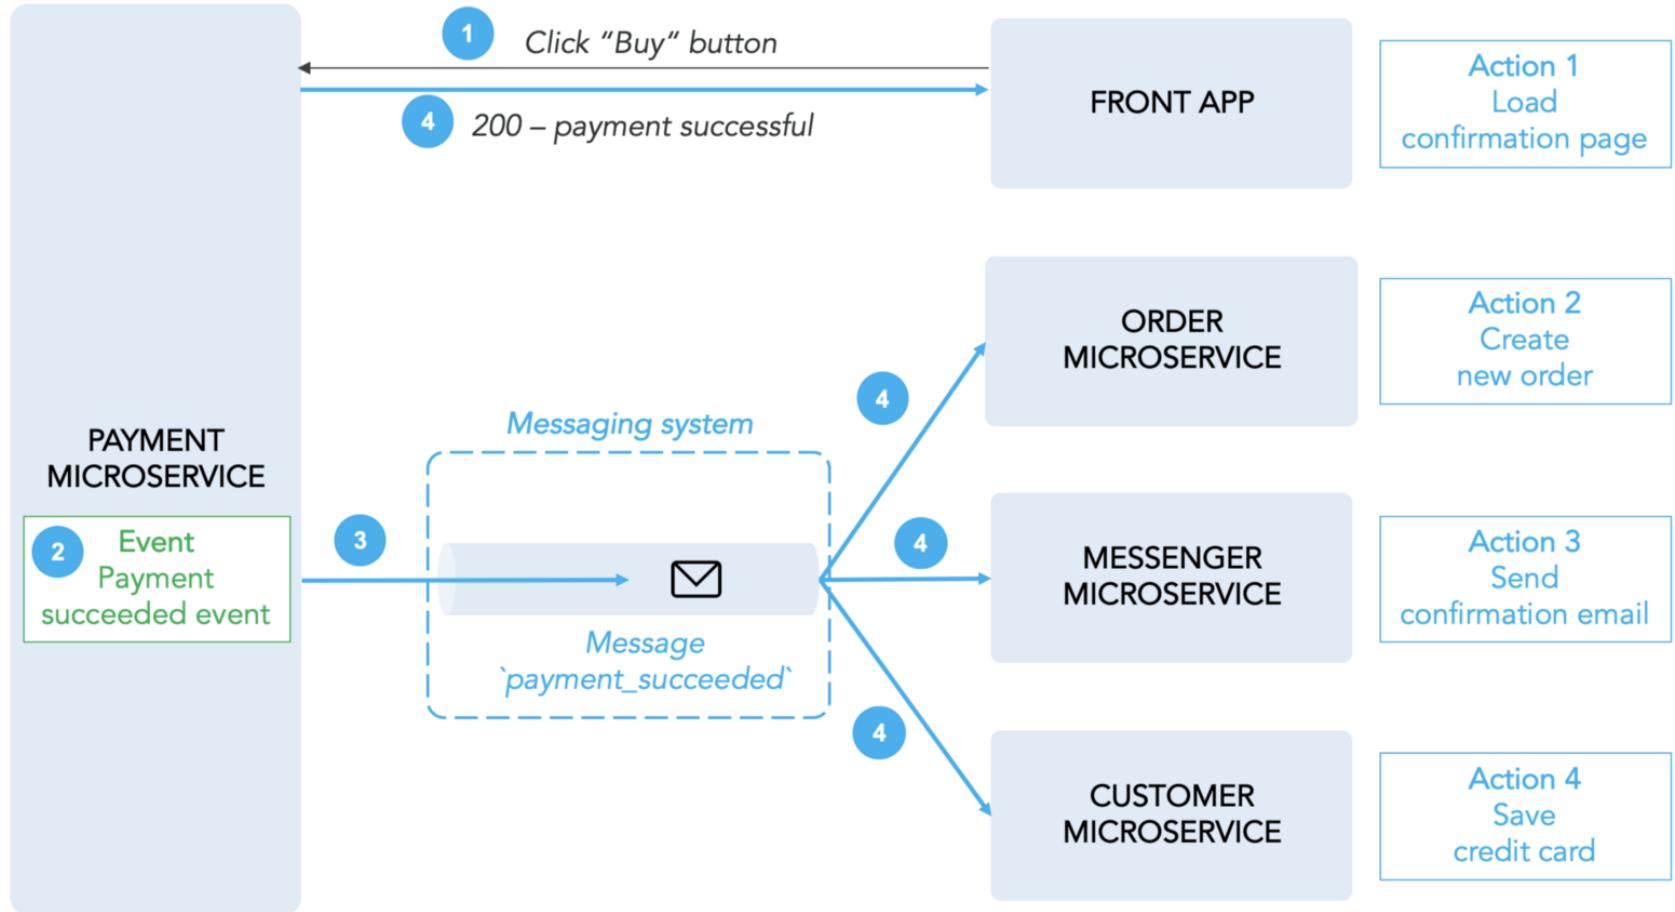
- ❖ Application development is currently based on concepts like microservices and serverless functions.
 - We now live in a cloud-native world, and these models are a natural fit for distributed, cloud-based environments.
 - But simply building and deploying services and functions isn't enough.
- ❖ On its own, a single microservice doesn't accomplish much.
 - We also need a way to wire up those components
 - i.e., to connect them so they can exchange data, forming a true application.
 - In SE, this is one of the most important architectural decisions to make.

Request-driven architecture



<https://blog.theodo.com/2019/08/event-driven-architectures-rabbitmq/>

Event-driven architecture



Event-driven architecture

- ❖ For controlling these services, various mechanisms were developed over the years, such as message queues and enterprise service buses (ESBs).
 - E.g. RabbitMQ, WSO2.
- ❖ More recent offerings, the concept of streaming data have also emerged.
 - E.g. Apache Kafka
- ❖ This latter category is growing, in part because streaming data is seen as a useful tool for implementing event-driven architecture
 - a software design pattern in which application data is modeled as streams of events, rather than as operations on static records.

What is an event?

- ❖ **Events** are things that happen, within a software system or, more broadly, during the operation of a business or other human process.
 - e.g., a sensor reports a temperature change, a user clicks their mouse, a customer deposits a check into a bank account.
- ❖ The concept of events in software systems closely aligns with how most of us think about our day-to-day lives.
 - Organizing around events makes it easier to develop business logic that accurately models real-world processes.
 - It helps reducing the number of one-to-one connections within a distributed system increasing the value of the microservices.
- ❖ An **event-driven architecture** allows generating, storing, accessing and reacting to these events.

Events vs. queries and commands

- ❖ **Queries** are a request to look something up
 - Unlike events or commands, queries are free of side effects; they leave the state of the system unchanged.
 - ❖ **Commands** are actions
 - Requests for some operation to be performed and will change the state of the system.
 - Synchronous and typically indicate completion.
-

	Behavior/state change	Includes a response
Command	Requested to happen	Maybe
Event	Just happened	Never
Query	None	Always

Event-driven patterns – notification

- ❖ **Event notification**
- ❖ A service sends events to notify other systems of a change in its domain.
 - For example, a user account service might send a notification event when a new login is created.
- ❖ What other systems choose to do with that information is largely up to them.
 - The service that issued the notification just carries on with its business.
- ❖ Notification events usually don't carry much data.
 - Resulting in a loosely coupled system with minimal network traffic spent on messaging.

Event-driven patterns – state transfer

- ❖ **Event-carried state transfer**
- ❖ A step up from simple notification, in this model the recipient of an **event** also receives the **data** it needs to perform further work.
 - E.g., the user account service might issue an event that includes a data packet containing the new user's login ID, full name, hashed password, and other pertinent details.
- ❖ This model can be appealing to developers familiar with RESTful interfaces.
 - But, depending on the complexity of the system, it can lead to a lot of data traffic on the network and data duplication in storage.

Event-driven patterns – sourcing

- ❖ **Event-sourcing**
- ❖ The goal of this model is to represent every change of state in a system as an event, each recorded in chronological order.
- ❖ In so doing, the **event stream itself becomes the principal source of truth for the system.**
 - E.g., it should be possible to “replay” a sequence of events to recreate the state of a SQL database at a given time.
- ❖ This model presents a lot of possibilities, but it can be challenging to get right.
 - particularly when events require participation from external systems.

Event-driven advantages (over REST)

❖ Asynchronous

- Allows resources to move to the next task once a unit of work is complete.
- Events are queued or buffered which prevents consumers from putting back pressure on producers or blocking them.

❖ Loose Coupling

- Services operate independently, without knowledge of other services, including their implementation details and transport protocol.
- Services under an event model can be updated, tested, and deployed independently and more easily.

❖ Easy Scaling

- Since the services are decoupled and typically perform only one task, tracking bottlenecks and scaling a service is easier.

❖ Recovery Support

- Can recover lost work by “replaying” events from the past.

<https://dzone.com/articles/best-practices-for-event-driven-microservice-archi>

Event-driven disadvantages

- ❖ They are easy to **over-engineer** by separating concerns that might be simpler when closely coupled.
 - they can require significant upfront investment, and often result in additional complexity, service contracts or schemas, polyglot build systems, and dependency graphs.
- ❖ **Complex data and transaction** management.
 - Typically do not support ACID transactions.
 - Systems must carefully handle inconsistent data between services, incompatible versions, duplicate events.
- ❖ Even with these drawbacks, ...
 - an event-driven architecture is usually the better choice for enterprise-level microservice systems.
 - The pros—scalable, loosely coupled, dev-ops friendly design—outweigh the cons.

Event-driven components

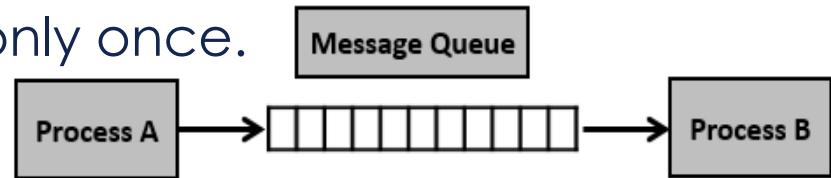
- ❖ Source/sink
 - In an E-D architecture, virtually everything generates events: devices, apps, user activity, etc. Many others consume them.
- ❖ Messaging system
 - This is the way events flow in the architecture, and the central element of the Event Notification pattern.
- ❖ Server-side applications
 - This is our code, e.g. web apps, microservices, functions, event processing apps, or rich data pipelines.
- ❖ Storage and analytics
 - An E-D architecture changes how we think about storage
 - e.g., if we do event sourcing, then we also record events.
- ❖ Event logging and monitoring
 - A downside of E-D is the lack of predictability.
 - The architecture must have a robust method for tracing event paths, monitoring app performance, and making troubleshooting tolerable.

Messaging systems

Messaging systems – models

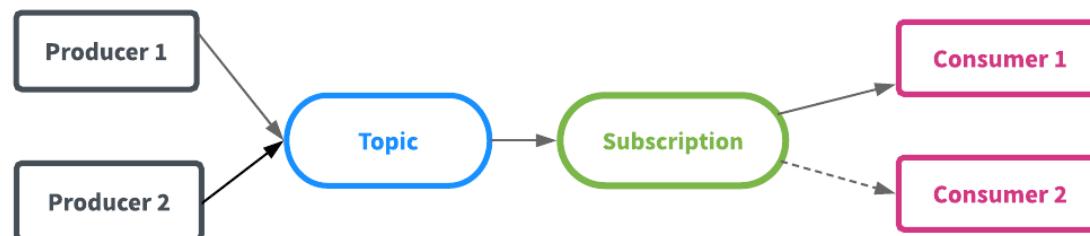
❖ Point to Point – Message queue

- Messages are sent to a queue to pre-defined receivers.
- One-to-one relationship between sender and consumer.
- Each message is consumed only once.



❖ Publish-Subscribe

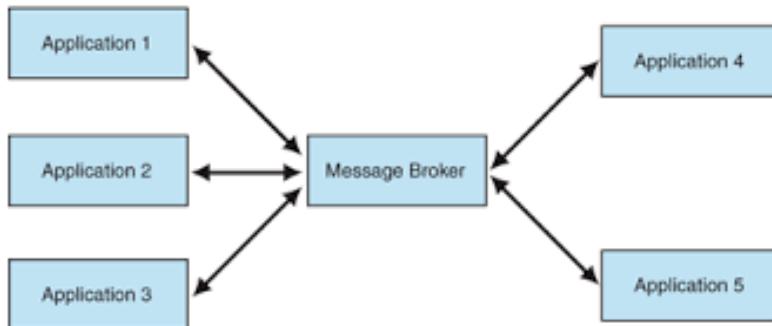
- Each message is published to a topic, and every application that subscribes to that topic gets a copy of all messages published to it.
- Message producers are also known as publishers and consumers are known as subscribers.



Messaging systems

- ❖ Managing the messages' flow
 - Messages are “put into” a source queue.
 - They are then “taken from” a destination queue.
 - How/Who/What moves a message from a source queue to a destination queue?

- ❖ Queue Manager / **Message Broker**
 - Function as message-queuing “relay” that interact with distributed applications.



Message broker

- ❖ Message Broker is built to extend MQ
 - it can understand the content of each message that it moves through the Broker.

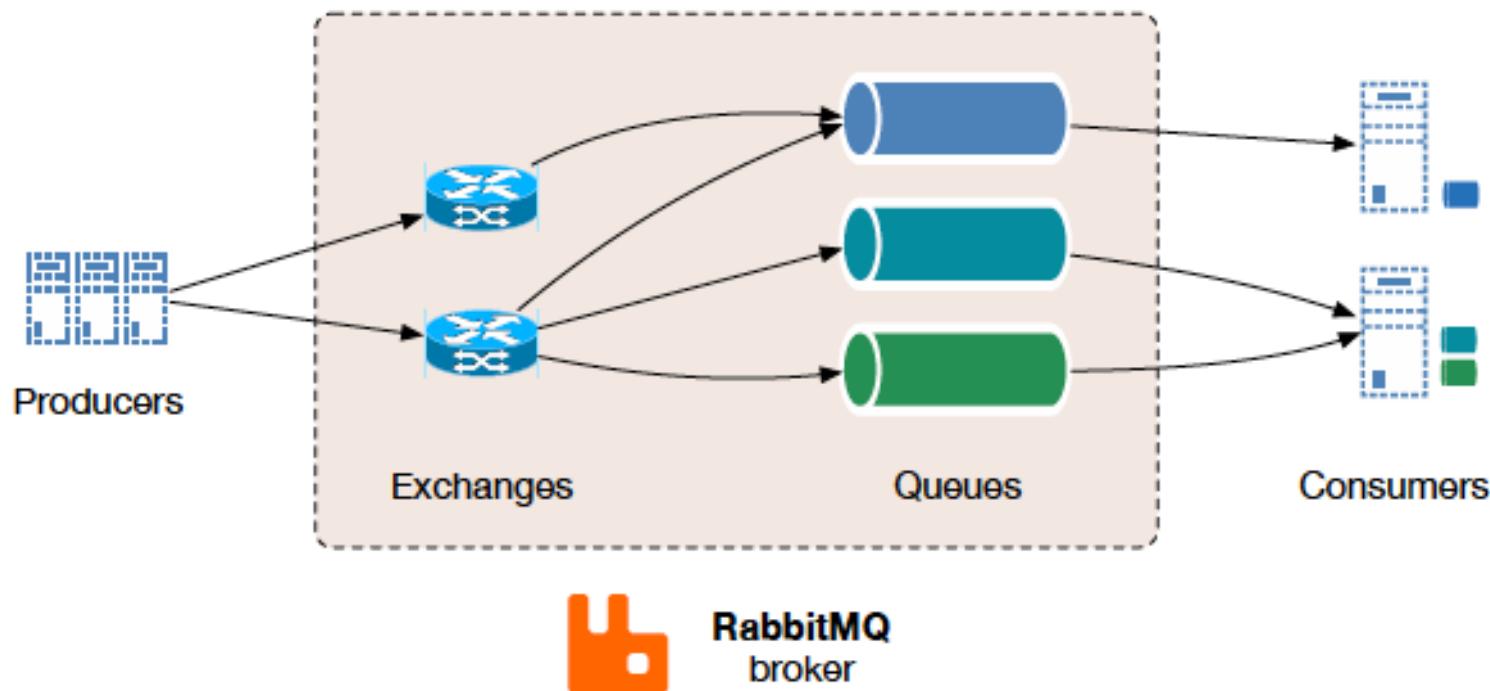
- ❖ Message Broker can do the following:
 - divide the publisher and consumer
 - store and route the messages between services
 - converts between different transport protocols
 - identifies and distributes business events from disparate sources

Message broker

- ❖ When a message broker is needed?
 - If we want to control data feeds, e.g., the number of registrations in a system.
 - When the task is to put data to several applications and avoid direct usage of their API.
 - When there is a need to complete processes in a defined order like a transactional system.
- ❖ There are many messaging tools..
 - E.g. Apache ActiveMQ, RabbitMQ
- ❖ .. and protocols
 - E.g. AMQP (Advanced Message Queuing Protocol), MQTT (MQ Telemetry Transport)

RabbitMQ

- ❖ RabbitMQ is one such open source message broker software that implements AMQP.



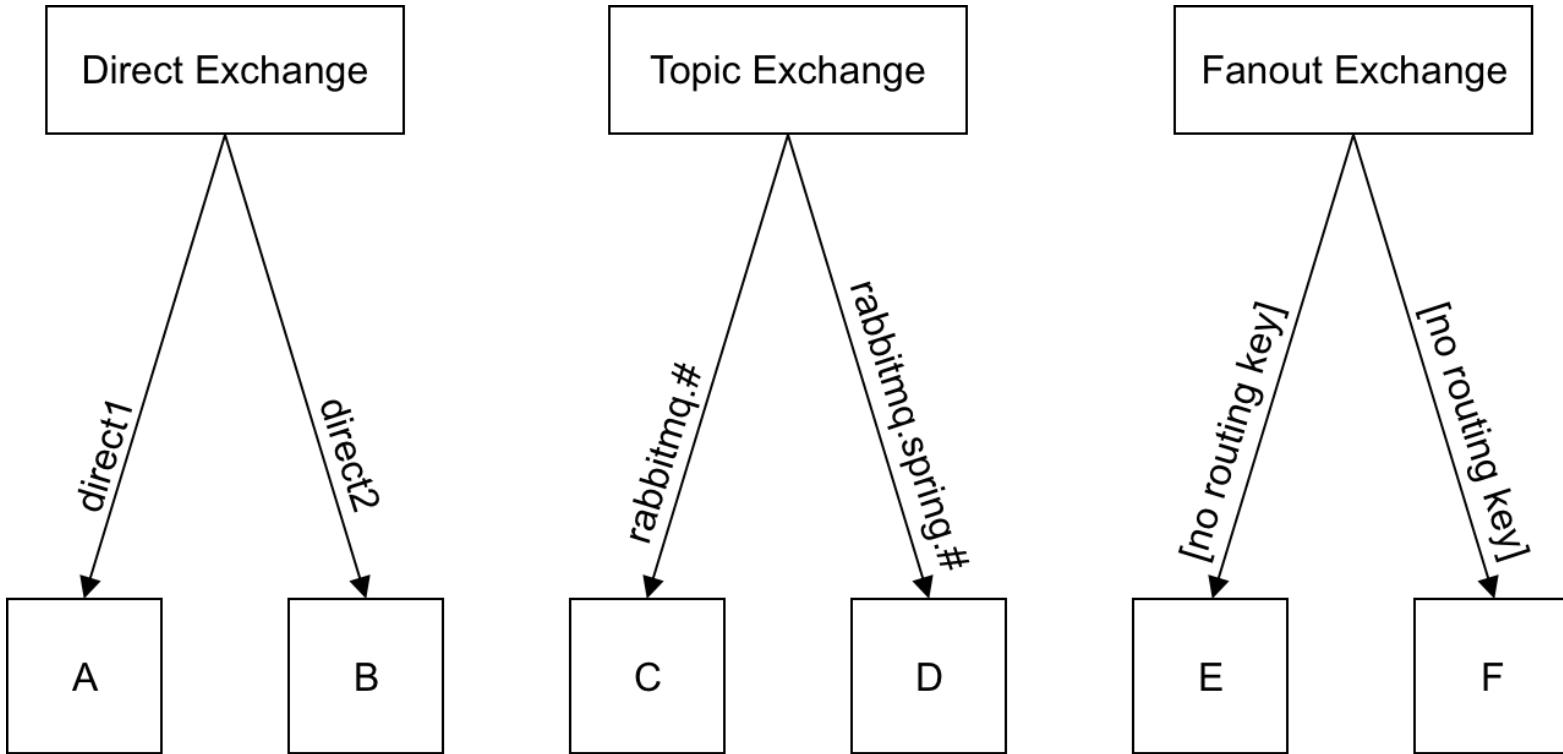
RabbitMQ – Main concepts

- ❖ **Messages** contains attributes (like headers in a request) and a payload (the message content).
- ❖ Messages are published to an entity, **exchange**, which distribute the messages to **queues** (or Topics).
- ❖ The rules for delivering the messages to the right queues are defined through
 - **bindings** (links between exchanges and queues), and
 - **routing keys** (a specific message attribute used for routing).
- ❖ Messages stored in queues are
 - delivered continually to subscribers, or
 - fetched by consumers on demand.

RabbitMQ – Exchange types

- ❖ **Direct Exchange** - It routes messages to a queue by matching routing key equal to binding key.
- ❖ **Fanout Exchange** - It ignores the routing key and sends message to all the available queues.
- ❖ **Topic Exchange** – It routes messages to multiple queues by a partial matching of a routing key. It uses patterns to match the routing and binding key.
- ❖ **Headers Exchange** – It uses message header instead of routing key.
- ❖ **Default(Nameless) Exchange** - It routes the message to queue name that exactly matches with the routing key.

RabbitMQ – Exchange types



Spring Boot – Messaging with RabbitMQ

- ❖ Set up the RabbitMQ Broker
- ❖ Spring Initializr
- ❖ E.g. RabbitMQ dependency

```
<!-- https://mvnrepository.com/artifact/org.springframework.boot/spring-boot-starter-amqp -->
<dependency>
    <groupId>org.springframework.boot</groupId>
    <artifactId>spring-boot-starter-amqp</artifactId>
    <version>2.2.1.RELEASE</version>
</dependency>
```

- ❖ <https://spring.io/guides/gs/messaging-rabbitmq/>

Event streaming systems

Apache Kafka

Stored records vs Messaging

ETL/Data Integration	Messaging
-	+
Batch	High Throughput
Expensive	Durable
Time Consuming	Persistent
	Maintains Order
	+
	Fast (Low Latency)
	-
	Difficult to Scale
	No Persistence
	Data Loss
	No Replay

<https://pt.slideshare.net/ConfluentInc/what-is-apache-kafka-and-what-is-an-event-streaming-platform>

The Event Streaming Paradigm



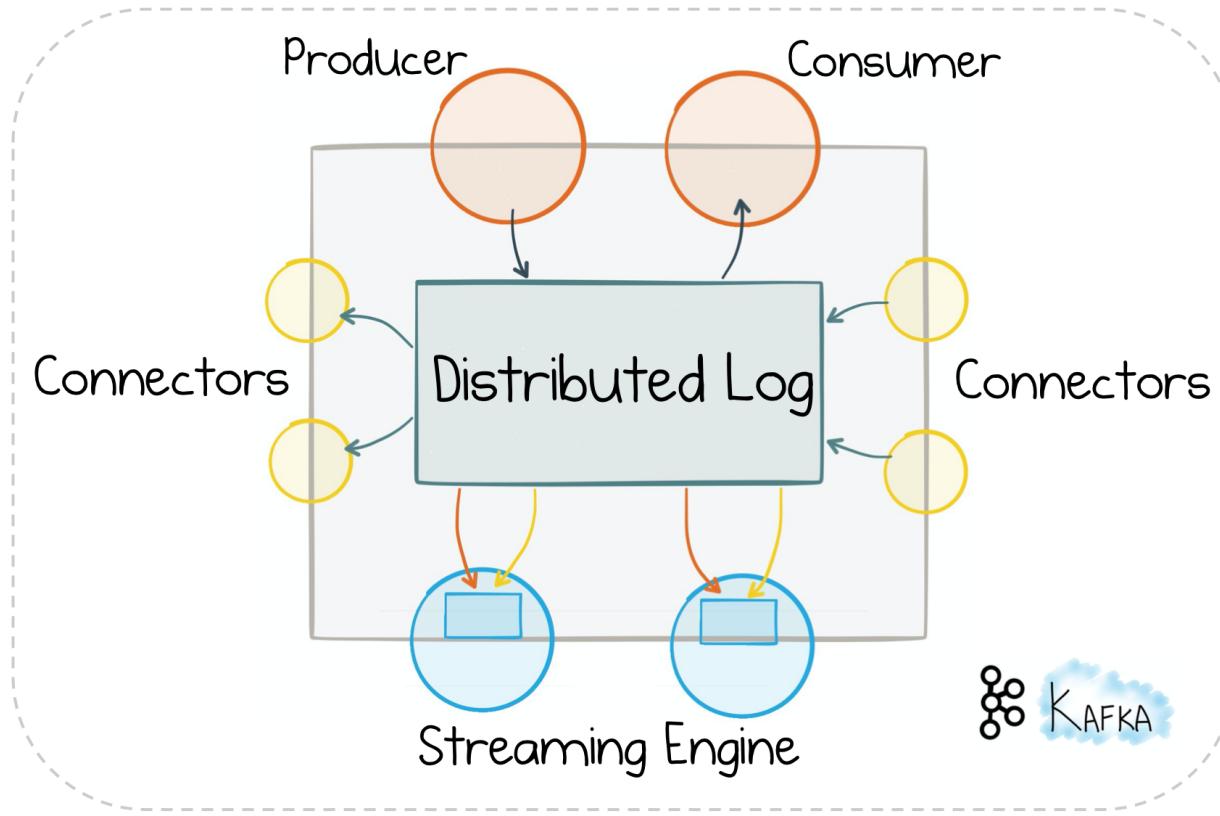
The Event Streaming Paradigm

To rethink data as not stored records or transient messages, but instead as a continually updating stream of events

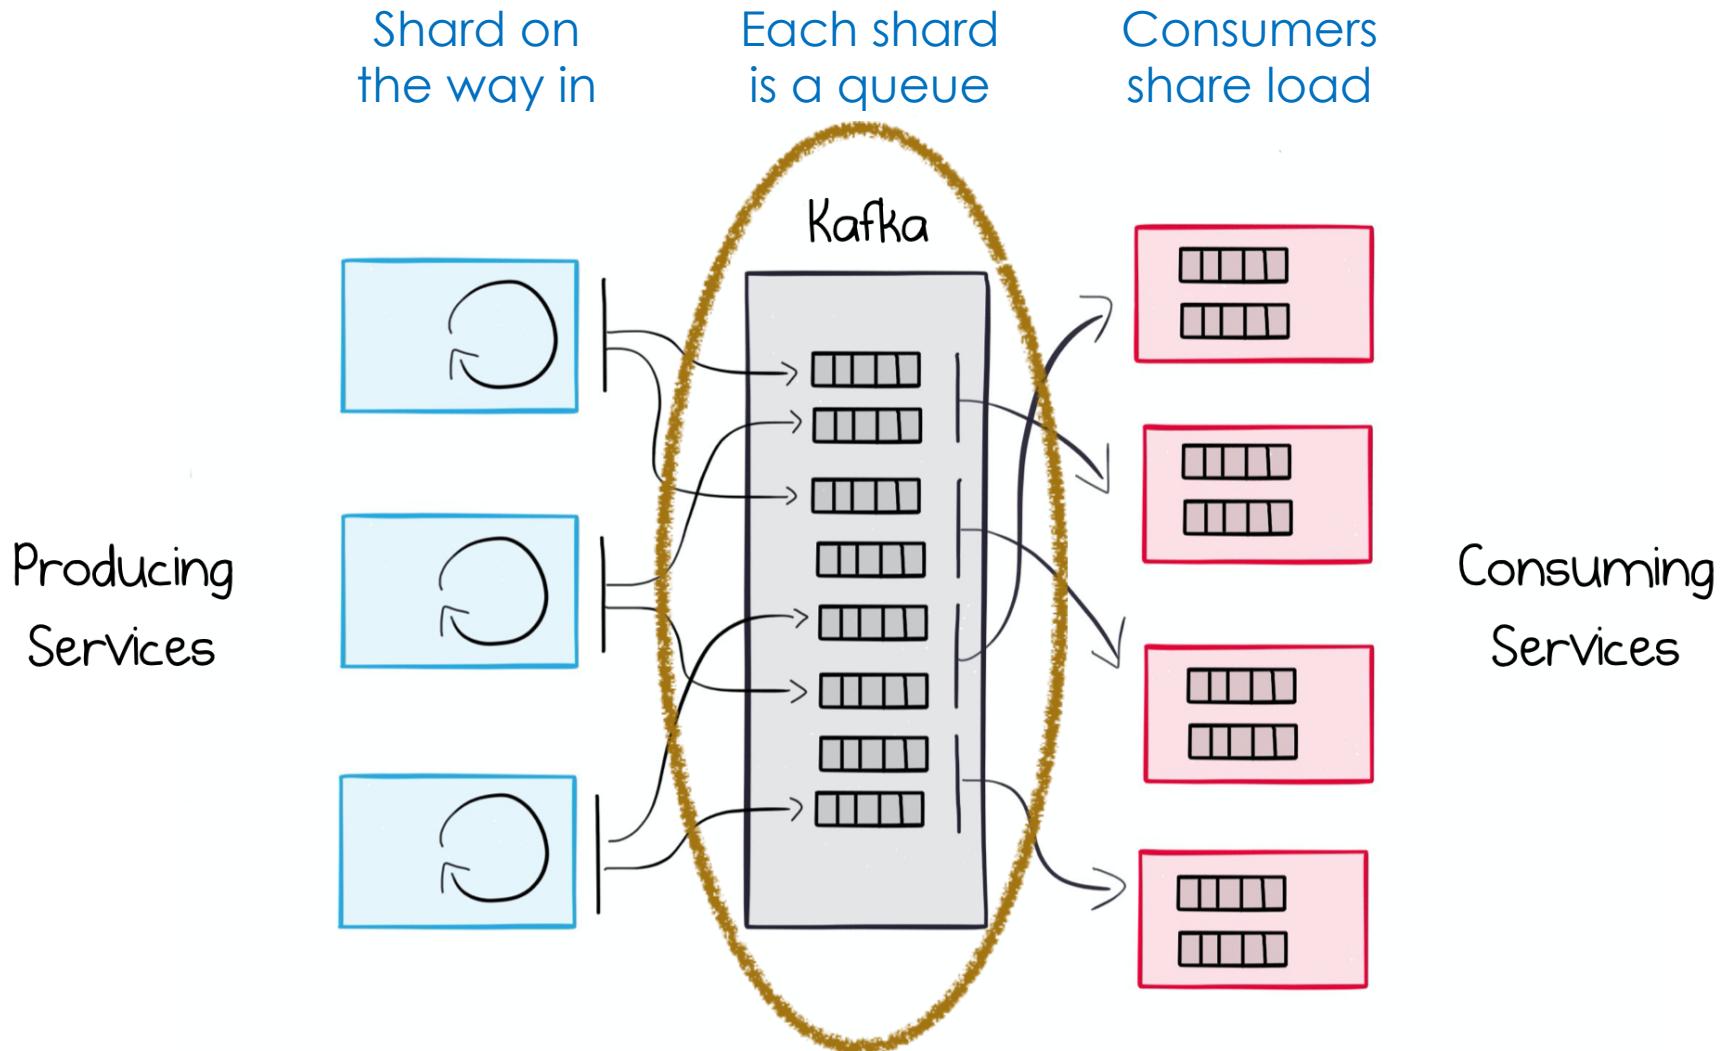
<https://pt.slideshare.net/ConfluentInc/what-is-apache-kafka-and-what-is-an-event-streaming-platform>

Apache Kafka

- ❖ Apache Kafka is made of distributed, immutable, append-only commit logs



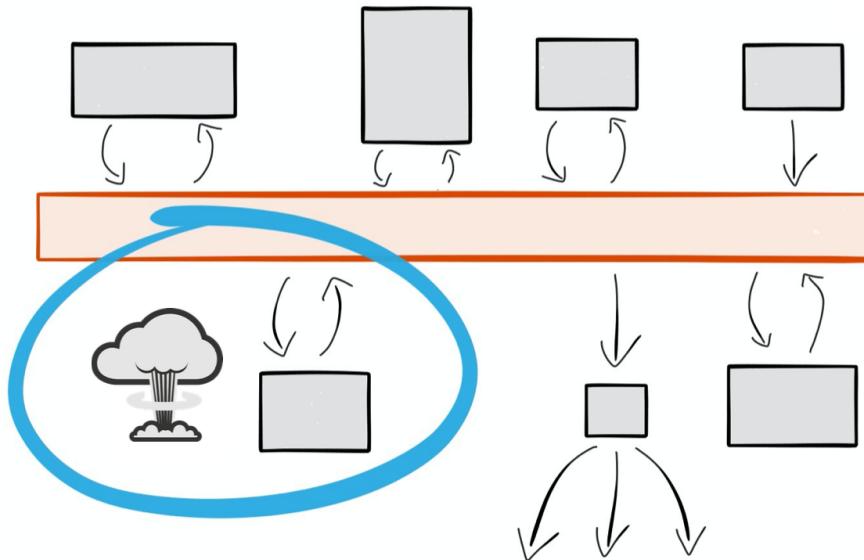
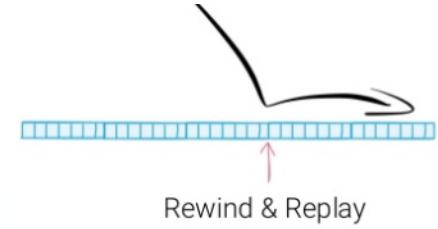
The Distributed Log



The Distributed Log

❖ The Service Backbone

- Scalable/Load Balanced, Fault Tolerant, Concurrent, Strongly Ordered, Stateful



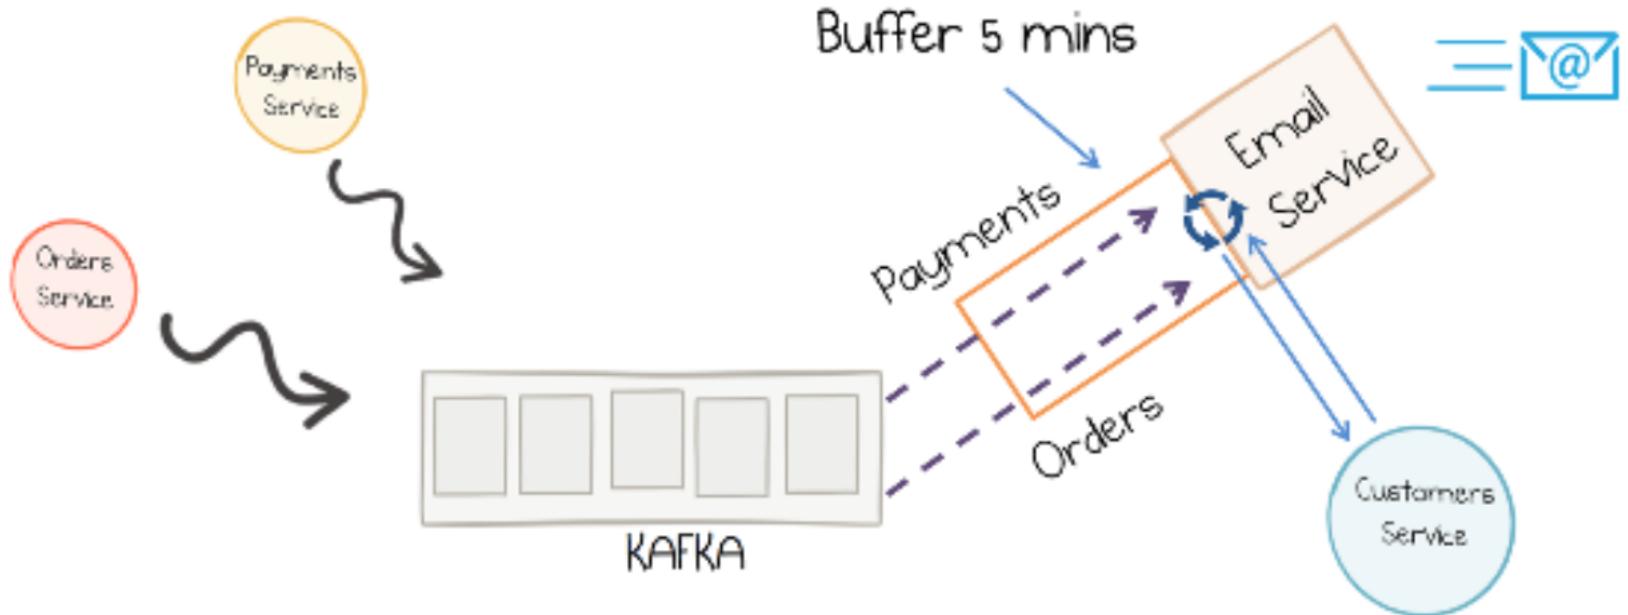
❖ A place to keep data-on-the-outside

Event-driven patterns

- ❖ Previously...
 - Event notification, Event-carried state transfer, Event-sourcing
- ❖ In **Event Sourcing**, events are a core element – the source of truth.
 - Being stored, immutably, in the order they were created in, the event log expresses exactly what the system did.
- ❖ **Command Query Response Segregation (CQRS)** is a natural progression from this.
 - Decoupling writing and reading operations.
 - As a simple example, we might write events to Kafka (write model), read them back, and then push them into a database (read model).

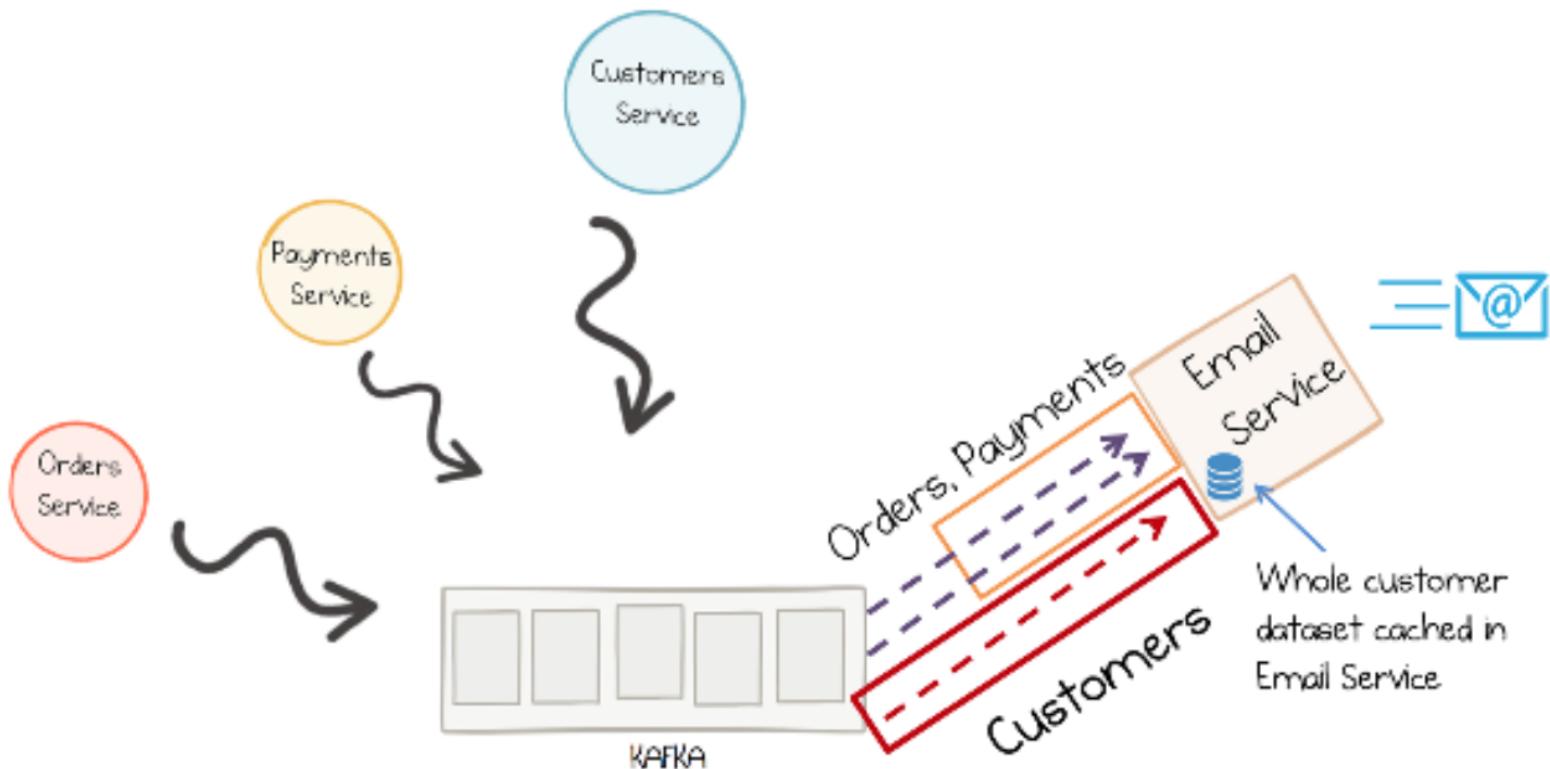
Stateless Streaming approach

- ❖ Example: A **stateless** streaming service that looks up reference data in another service at runtime



Stateful Streaming approach

- ❖ Example: A **stateful** streaming service that replicates the Customers topic into a local table, held inside the Kafka Streams API.



Stateful Streaming approach

❖ Disadvantages:

- The service is now stateful, meaning for an instance of the email service to operate it needs the relevant customer data to be present.
- This means, in the worst case, loading the full dataset on startup.

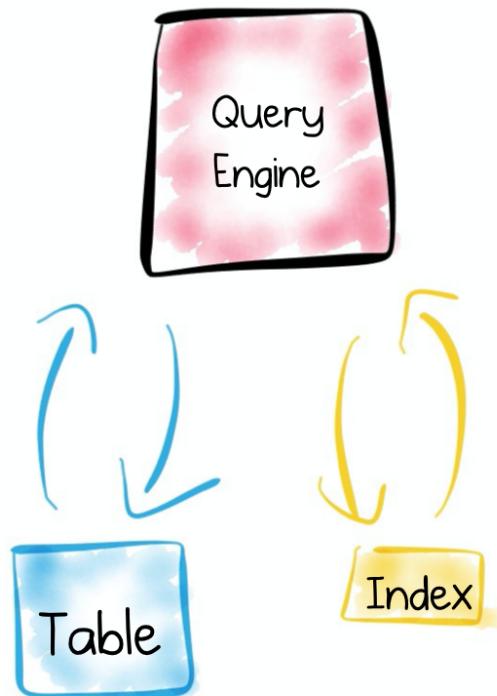
❖ Advantages:

- The service is no longer dependent on the worst-case performance or liveness of the customer service.
- The service can process events faster, as each operation is executed without making a network call.
- The service is free to perform more data-centric operations on the data it holds.

Stateful Streaming approach

- ❖ Being stateful comes with some challenges:
 - when a new node starts, it must load all stateful components (i.e., state stores).
- ❖ Kafka Streams, for instance, provides three mechanisms to simplify stateful:
 - It uses a technique called **standby replicas**, which ensure that for every table or state store on one node, there is a replica kept up to date on another.
 - **Disk checkpoints** are created periodically so that, should a node fail and restart, it can load its previous checkpoint.
 - Finally, **compacted topics** are used to keep the dataset as small as possible. This acts to reduce the load time for a complete rebuild should one be necessary.

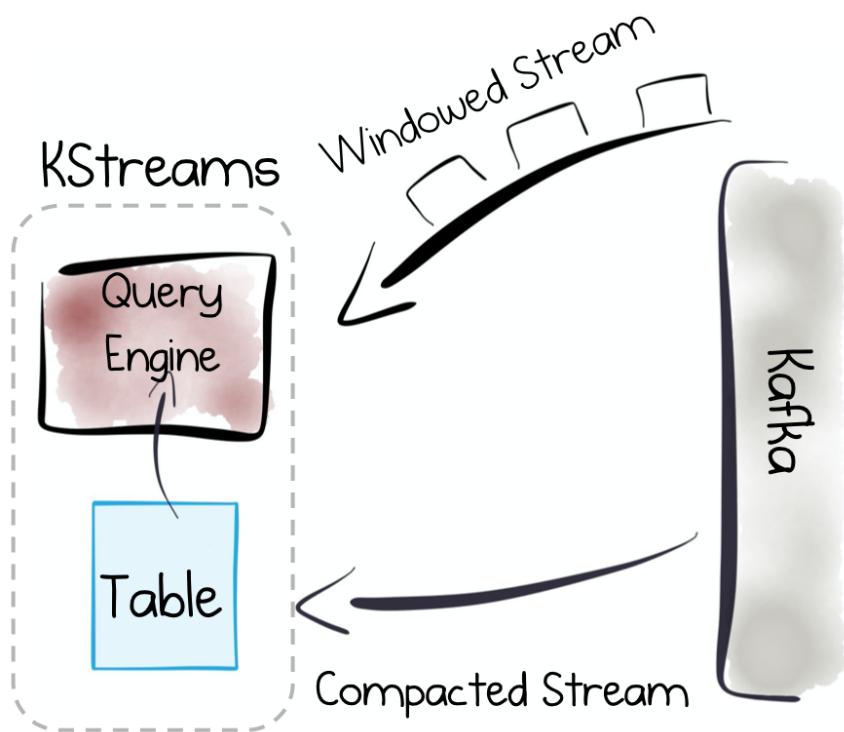
A Stateful Stream Processing



Database

Finite source

VS

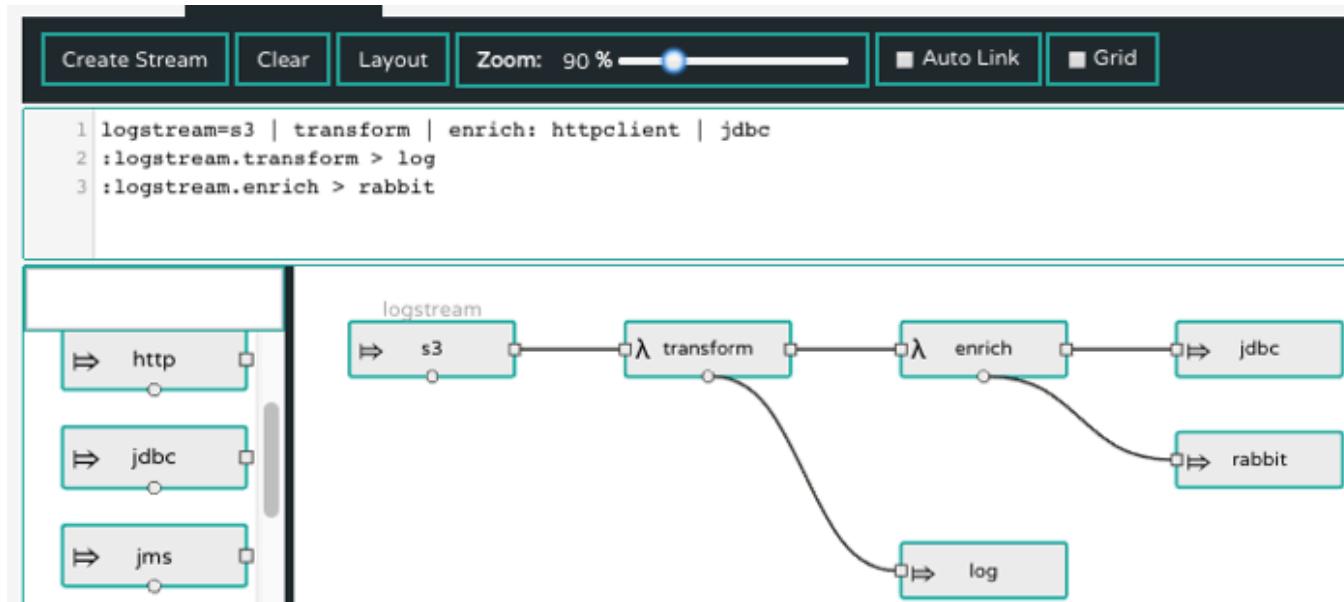


Stateful Stream Processor

Infinite & Finite source

Streaming in Spring

- ❖ **Spring Cloud Stream** helps fully abstract code from the underlying messaging engine.
- ❖ **Spring Cloud Data Flow** allow to create and orchestrate data pipelines, e.g. data ingest, real-time analytics, and data import/export.



Summary

- ❖ Event-driven architecture is gaining in popularity, and with good reason.
 - From a technical perspective, it provides an effective method of wiring up microservices.
 - The interest in serverless functions - such as AWS Lambda, Azure Functions, or Knative - is growing, and these are inherently event-based.
 - Moreover, when coupled with modern streaming data tools like Apache Kafka, event-driven architectures become more versatile, resilient, and reliable than with earlier messaging methods.
- ❖ But perhaps the most important “feature” of the event-driven pattern is that it models how businesses operate in the real world.